



The HDF Group



Why HDFhub?

JupyterHub for HDF

Why Jupyter Notebooks?

- Step by step notebook for a user to understand and replicate
- Switch between Python, BASH, etc in one GUI, sharing variables between languages along with other “magics”
- Opportunity to inject markdown cells to explain in plain English what is happening in each chunk of code.
- Images, video are embeddable, so you could have a short video that explains a section of code or the slides for a presentation



All of this is really cool, but...

- A notebook is dependent on a Notebook or Lab server.
- The application is usually run locally, and on multiple operating systems, so syntax differences exist
- This results in a situation where the notebook might not work in the environment that someone else is using.



A notebook is not a notebook

- Homegrown
 - Specific focus to environment
 - unsupported
 - contributable
- Open Source
 - most commonly used packages
 - supported
 - contributable
- Enterprise
 - most comprehensive
 - Supported
 - Contributable



Why JupyterHub?

“But the plans were on display . . .”

“On display? I eventually had to go down to the cellar to find them.”

“That’s the display department.”

“With a torch.”

“Ah, well the lights had probably gone.”

“So had the stairs.”

“But look, you found the notice, didn’t you?”

“Yes,” said Arthur, “yes I did. It was on display in the bottom of a locked filing cabinet stuck in a disused lavatory with a sign on the door saying Beware of the Leopard.”

—Douglas Adams, The Hitchhiker’s Guide to the Galaxy

A Notebook is not enough

- The addition of JupyterHub allows us to control the user's environment and data
- This results in redoability not just reproducibility

Science is show me not trust me. - Philip Stark

Other Options to improve notebooks



nbviewer is integrated with Github

- Renders a static version of the last run, committed state of a notebook in the branch you are looking at
- Doesn't allow interactivity
- Doesn't share the environment

Create a function to select the organization the metadata comes from

```
In [3]: def OrganizationChoices(organization):  
        global OrganizationChoice  
        global Organization  
        Organization=organization  
        print("Organization of the collection is", Organization)
```

Create a dropdown using the Organizations list and the organization selector function. This sets the Organization variable.

```
In [4]: interactive(OrganizationChoices, organization=Organizations)
```

Create a list of collections in the organization directory selected in the dropdown above

```
In [5]: Collections = []  
        for (dirpath, dirnames, filenames) in walk(os.path.join('../collection', Organization)):  
            Collections.extend(dirnames)  
            break  
        Collections
```

```
Out[5]: ['GES_DISC', 'GHRC', 'LARC', 'NSIDC', 'PODAAC']
```




Projects like Binder are promising

Turn a GitHub repo into a collection of interactive notebooks

Have a repository full of Jupyter notebooks? With Binder, open those notebooks in an executable environment, making your code immediately reproducible by anyone, anywhere.

Build and launch a repository

GitHub repo or URL

`https://github.com/scgordon/MILE2`

Git branch, tag, or commit

master

Path to a notebook file (Optional)

Path from repo root to notebook file (optional)

launch

Waiting Building

Build logs [show](#)



But it is beta...

Secure <https://hub.binder-beta.omgwtf.in/user/91ecd38e-2d9e-461f-b7d8-11d56b758c93/notebooks/notebook/Metadata2Data.ipynb>

jupyter Metadata2Data Last Checkpoint: 5 minutes ago (unsaved changes) Control Panel Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

The notebook utilizes Bash and Python with the default packages contained in the default build of Anaconda with Python 3.6.

Saxon, Java, and XSLT form the evaluation web service on a NCEAS virtual machine.

This CSV contains a row for each concept that is found, so some elements may fulfill multiple concepts. A good example of this are the concepts Keyword and Place Keyword. Every Place Keyword is also a Keyword, so the row would repeat with a different Concept name. It also contains a row for each undefined node that contains text, marking these rows with an Unknown in the Concept column.

This data can be used in a variety of analyses including RAD and QuickE as well as Concept Verticals. It can also be used to teach the system dialect definitions for concepts that are currently unknown by exposing all of the content at undefined nodes.

First we need to call all of the libraries we need to perform in our metadata wrangle

```
In [1]: import pandas as pd
import os
from os import walk
import shutil
from ipywidgets import *
import ipywidgets as widgets
import requests
from contextlib import closing
import csv
import io
```

```
-----
ImportError                                Traceback (most recent call last)
<ipython-input-1-fe2d6e2b09b0> in <module>()
----> 1 import pandas as pd
      2 import os
      3 from os import walk
      4 import shutil
      5 from ipywidgets import *
```

ImportError: No module named 'pandas'

Building HDFhub



Build a JupyterHub

You can build your own from scratch

- Can change the authenticator to the website's
- Can choose spawner
- Can build the environment
- Can use Notebook or Lab for the GUI

Or just grab a docker container and get started
jupyterhub-deploy-docker

jupyterhub-deploy-docker provides a reference deployment of [JupyterHub](#), a multi-user [Jupyter Notebook](#) environment, on a **single host** using [Docker](#).

Possible **use cases** include:

- Creating a JupyterHub demo environment that you can spin up relatively quickly.
- Providing a multi-user Jupyter Notebook environment for small classes, teams, or departments.

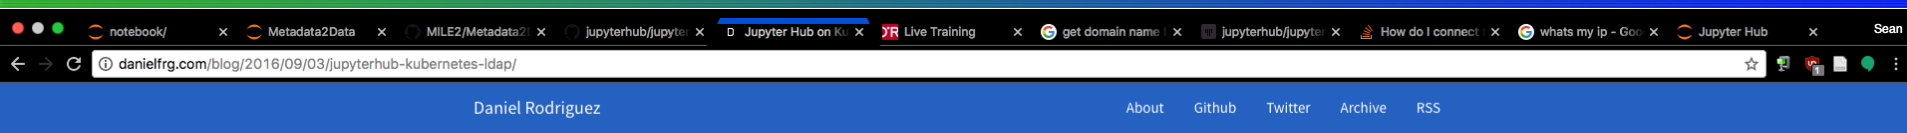


Kubernetes

- Yuvi Panda: Jupyter Hub can be made safer using Kubernetes bc Kubernetes has a bunch of ways to limit impact. No root, remove SMTP access, limit number of sockets. Google or other cloud handles persistence of data if desired
- Can create multiple JupyterHubs with different environments
- Treats the vm like cattle not pets
- Likely to see the most active development and optimization as Yuvi is a core developer for kubespawner as well as Jupyter and Data8



Kubernetes with Authentication and Storage



Jupyter Hub on Kubernetes with LDAP

September 03, 2016

In this post I am going to show some initial work I did in the last day to deploy Jupyter Hub in Kubernetes with user auth based on LDAP.

I wasn't that much work considering that Jupyter Hub already had support for LDAP user auth and the modularity they have is amazing. It was quite straight forward to write a new spawner based on the existing one on the [Jupyter Hub github org](#).

All the code in this post is at [danielfrg/jupyterhub-kubernetes_spawner](#). Specifically in the `examples` directory.

I consider this a nice example of a more production ready Jupyter Hub deployment being based on LDAP that for good or bad is everywhere and Kubernetes to deployment of the single user notebook servers.

Something I don't talk is SSL certificates because its very well supported on Jupyter Hub and I think there is enough information about them on the [Jupyter Hub docs](#) and online. And yes, I was lazy to



Jupyter Hub on Kubernetes Part II: NFS

September 10, 2016

Second part of my JupyterHub deployment in Kubernetes experiment be sure to read [Part I](#).

Last time we got JupyterHub authenticating to LDAP and creating the single user notebooks in Kubernetes containers. As I mentioned in that post one big problem with that deployment was that the notebook files are gone when the pod is deleted so this time I add an NFS volume to the JupyterHub single user containers to persist the notebook data.

Also improved a little bit the deployment and code so to its no longer needed to build a custom image you can just pull two images from my docker hub registry and configure them using a [Kubernetes ConfigMap](#).

All the code in this post is at [danielfrg/jupyterhub-kubernetes_spawner](#). Specifically in the `example` directory.

NFS



VM based tutorial got me a lot further

<http://localhost:8080/user/vagrant/tree/MILE2>



Gitter for all the technical questions

- The majority of Jupyter developers are active in this Gitter Channel:

<https://gitter.im/jupyterhub/jupyterhub>

The screenshot shows the Gitter chat interface for the #jupyterhub channel. The interface is divided into three main sections: a left sidebar for navigation, a central chat area, and a right sidebar for user and activity information.

Left Sidebar: Contains navigation links for COMMUNITIES, JUPYTERHUB, Home, ALL ROOMS, and a list of rooms including jupyterhub/jupyterhub and binder.

Central Chat Area: Displays a list of messages. The messages are from users like @ldacey, @minrk, @willingc, @cmoscardi, @yuvi, @dlsun, and @Winterflower, discussing technical issues related to JupyterHub deployment and authentication.

Right Sidebar: Contains a section for PEOPLE (653 PEOPLE) and a section for ACTIVITY. The ACTIVITY section shows recent actions like 'synchronize #1408', 'unlabeled #1361', 'labeled #1409', 'closed #1410', 'commented #1410', 'opened #1410', 'opened #1409', 'commented #1361', and 'commented #1372'.



Examples of JupyterHub in the wild

- Data8 <http://data8.org/>
<https://github.com/data-8/jupyterhub-deploy>
- Harvard <https://github.com/harvard/cloudJHub>
- Quantopian Research Platform
<https://libraries.io/github/quantopian/jupyterhub>
- Lawrence Berkeley Labs <https://goo.gl/GBQ5gS>
- Enterprise solutions Anaconda, Immuta, DataScience offering free trials
- O'Reilly using Hub and Notebooks for webinars now

Why HDFhub for website engagement and beyond



Hub as community engagement

- Interactive with a lower barrier to entry
- People can easily and quickly build on each other's ideas
- This makes HDF more accessible in the perception of our user base as a format and support
- Likely to stimulate conversation on the forum since they are also on the website
- Identify development opportunities to expand HDF capabilities and tools by interacting with current and future users



No more workshop software carpentry

- Allows us to focus on the why immediately
- Only a connected web browser is needed
- Provides an understood GUI with powerful documentation capabilities to drive interactions with products like HSDS or the metadata evaluator
- Gives tangible results to users
- Easily tweakable to particular use case



Larger set of potential presenters

- This means we can reach out on more fronts at any conference/workshop/etc sending fewer participants.
- People more directly involved in identifying potential revenue can show all off our work easily rather than just tell about it with potential customers.
- Allows us all to play to our strengths rather than trying to fill all roles.



Where do we go now?

- Jupyter Notebooks are necessary but not sufficient
- The right environment is needed as well
- What we could do
 - Use the free Binder beta with all it's warts while it's free and hope it works when it's needed
 - Get our own Binder instance running so developers can keep their own work and requirements document up to date
 - Create a JupyterHub that gives each website user and employee persistent a notebook/lab server to attract website users and identify needs and trends while engaging with a broader audience
 - Create a JupyterHub that gives each website user the ability to read and execute featured Notebooks in a Notebook/Lab environment but no persistence and create persistent cloud environments for employees



Binder

- This presentation has been hosted on Binder.
- It's part of the repository I used to give people a user interface to our metadata evaluation for improvement at the ESIP Summer meeting this year